

A taxonomy along three independent axes

Classify the error, the evidence used to detect it, and the point of intervention.

A What is the reference?

World factuality

Incorrect factual claims
Reference: external world
e.g., TruthfulQA; FActScore

Source faithfulness

Contradicted or unsupported
Reference: supplied evidence
e.g., RAGTruth; FaithBench

These criteria can disagree.

Unsupported is not always false;
source-consistent can be wrong.

B What evidence is used?

External / source evidence

NLI / QA / atomic claims
AlignScore; VeriScore

Sampled-text consistency

Repeated / cross-model outputs
SelfCheckGPT; semantic entropy

Model-internal signals

Probabilities / states / attention
FOCUS; Lookback Lens

C Where to intervene?

Training and alignment

FaithDial; Self-Alignment

Retrieval and prompting

Self-RAG; Astute RAG

Decoding and intervention

CAD; DoLa; TruthX

Revision and abstention

CoVe; CRITIC; SEAL

Cross-cutting comparisons: granularity · model access · supervision · cost · coverage · robustness

Original synthesis. A method may occupy multiple categories; the axes do not assert causal relationships.