

Judgement Hallucination Examples

Source: library record

Northbridge Library opened in 2020.

It holds 12,000 books.

The archive is on the second floor.



Is there a cafe in the library?



Yes, there is a cafe in the library.



Does it hold 20,000 books?



Yes, it holds 20,000 books.



Is the archive on the first floor?



Yes, the archive is on the first floor.

Description Hallucination Examples

Source: delivery note

200 books arrived Monday.

Sorting: Tuesday. Cataloguing:

Wednesday.



Summarize the delivery note.



On **Tuesday**, the library received **500 books**.

The books came **from a university**.

Cataloguing happened before **sorting**.

Unsupported does not necessarily mean false in the world; the source does not establish it.

Figure 1: Source-faithfulness hallucinations in constructed QA and summarization examples, not measured model outputs. Green marks unsupported additions; magenta and blue mark value and relation contradictions. Highlights are illustrative, not an exhaustive taxonomy; world factuality is a separate criterion.




Figure 2: The structure of this survey with representative references. Each green leaf includes four to five studies. Evaluation criteria are detailed in Fig. 3 and Table 1, and intervention mechanisms in Fig. 4. Hybrid studies may appear in more than one section.




Figure 3: Taxonomy of evaluation targets and criteria with benchmark references. Gen evaluates generated content; Dis evaluates detectors against human labels. Criteria are not shared protocols; Table 1 gives benchmark-specific details. The same benchmark may support both targets.




Figure 4: Potential failure sources and representative mitigation methods in text LLMs. Retrieval is optional. Stage associations summarize the literature; they do not establish one-to-one causal effects or guaranteed fixes. Evaluate outcomes using Fig. 3 and Table 1.

Table 1: Hallucination evaluation benchmarks for text LLMs. Dis: hallucination discrimination; Gen: generated-content evaluation; MC: multiple choice. Size refers to the stated corpus unit, not a common test-set size.

Benchmark	Evaluation	Size	Hallucination focus	Metrics
TruthfulQA [Lin <i>et al.</i> , 2022]	Gen / MC	817 questions	World factuality	Truthfulness; informativeness
HaluEval [Li <i>et al.</i> , 2023]	Dis	35,000 samples	Mixed-task hallucination	Recognition accuracy
ANAH [Ji <i>et al.</i> , 2024]	Dis	~12k sentences	QA support / contradiction	Type accuracy; F1
RAGTruth [Niu <i>et al.</i> , 2024]	Dis	17,790 responses	RAG source faithfulness	Response / span P, R, F1
ExpertQA [Malaviya <i>et al.</i> , 2024]	Gen	2,177 questions	Correctness & attribution	Human factuality / support
TofuEval [Tang <i>et al.</i> , 2024]	Dis / Gen	1,479 summaries*	Dialogue-summary consistency	Balanced accuracy; error rate
HALoGEN [Ravichander <i>et al.</i> , 2025]	Gen	10,923 prompts	Domain-specific factuality	Atomic hallucination rate
FaithBench [Bao <i>et al.</i> , 2025]	Dis / Gen	750 summaries	Summary source faithfulness	Balanced accuracy; macro F1
FactBench [Fatahi Bayat <i>et al.</i> , 2025]	Gen	1,000 prompts	Verifiable factual claims	Hallucination score; precision

* Retained count reported in the paper. Dataset sizes count different units; these rows are not a pooled leaderboard.
HALoGEN size is for the complete suite; only text-relevant tasks fall within this survey. Paper snapshots: 2022-2025.