

Hallucination in Large Language Models: A Survey of Detection, Evaluation, and Mitigation

Xincheng Li
University of Auckland

Abstract

Large language models can produce fluent assertions that are false or unsupported by their evidence. Research spans detectors, benchmarks, and interventions, but their results often concern different correctness targets. This survey organizes text-based research from 2022–2026 around error definitions, detection evidence, and useful mitigation. We distinguish world factuality from source faithfulness, compare external-evidence, sampling-based, and internal-signal detectors, and connect evaluation to training, retrieval, decoding, and revision. Factual precision alone cannot establish progress: verifier dependence, selective answering, evidence quality, and distribution shift change the interpretation of gains. We develop a comparative framework pairing correctness with information coverage, independent verification, and cost, and propose tests of generalization and reliability. The companion literature collection and supporting materials are available at <https://github.com/lixincheng-xcl/LLM-Hallucination-Survey>.

1 Introduction

A fluent answer is not necessarily justified. In question answering (QA), a model may repeat a misconception; in summarization, it may add a plausible detail absent from the source. Both are called hallucinations, although their correctness criteria differ: source-unsupported content can be factually correct (Lin et al., 2022; Cao et al., 2022). Collapsing these settings obscures what a system improves.

Figure 1 makes this distinction operational. The invented library record supports 12,000 books and a second-floor archive. The cafe assertion lacks support; the changed count and location contradict the record. The delivery summary similarly alters quantities and event order while inventing a donor.




Figure 1: Constructed source-faithfulness errors. Green adds unsupported information; pink changes attributes; blue changes relations. The cafe claim is unsupported, not necessarily false in the world. Layout adapted from Liu et al. (2024a); examples are original, not measured outputs.

Verification must therefore specify both the claim and its reference evidence.

The field now spans claim verification, uncertainty estimation, and interventions throughout generation (Wang et al., 2024). Their requirements differ: retrieval, repeated answers, or access to model internals. Lower error rates can also reflect shorter answers or more refusals. We organize this diversity around three questions: **RQ1**, what reference frame defines an error; **RQ2**, what evidence and access enable detection; and **RQ3**, when does mitigation improve useful generation?

Scope and organization. This targeted 2022–2026 survey emphasizes ACL-family work on text-based QA, long-form generation, summarization, dialogue, and retrieval-augmented generation (RAG); it excludes multimodal perception and general reasoning. Figure 2 maps concepts (§2), detectors (§3), evaluation (§4), interventions (§5), and critical synthesis (§6). Appendix A documents selection and reading depth. We report no new experiments or exhaustive systematic review.

2 Concepts and Scope

2.1 Reference frames

World factuality concerns agreement with externally established facts; *source faithfulness* concerns support by supplied evidence. Neither entails the other: a response can repeat a false source faithfully or add a true but unlicensed detail (Cao et al., 2022; Adlakha et al., 2024). BEGIN evaluates attribution in grounded dialogue under this distinction (Dziri et al., 2022b).

Source-based labels should distinguish *supported*, *contradicted*, and *not established*, retaining the last separately from world falsity. World verification additionally needs a specified evidence source, date, and unresolved-claim policy. These are operational choices, not a universal definition. HalluLens revisits definitions and benchmark construction (Bang et al., 2025); comparisons must map study-specific labels before comparing scores.

2.2 Error relations and development

Entity, attribute, and relation errors are orthogonal to reference frames: a wrong date can contradict either a document or a world fact. Annotation units also differ. ANAH provides sentence-level analysis, while HaluEval 2.0 extends empirical factuality analysis across domains (Ji et al., 2024; Li et al., 2024).

Research moved from distinct task-specific correctness targets in 2022 toward atomic verification and black-box self-checking in 2023 (Min et al., 2023; Manakul et al., 2023), followed by contextual monitoring and factuality alignment. Recent studies scrutinize judge reliability, information coverage, and temporal validity (Godbole and Jia, 2025; Samarinas et al., 2025; Jiang et al., 2026). This trajectory does not imply that newer methods supersede earlier ones.

We separate three axes: reference frame, detection evidence, and intervention point. A benchmark evaluates a detector; a detector can also guide decoding. Potential failure points include unfamiliar supervision, knowledge access, retrieval, and evidence use. Fine-tuning and long-context studies motivate these possibilities without proving a complete causal decomposition (Gekhman et al., 2024; Liu et al., 2024c).

Across these axes, we compare evidence access, generator signals, supervision, and computational cost. This separates a method’s decision criterion from its deployment requirements and supplies

common dimensions for the synthesis in Section 6.

3 Detection: Evidence and Access

Detection estimates response, sentence, or span errors under a specified reference. Figure 2 groups its evidence requirements: output access does not imply repeated samples, probabilities, or attention maps.

3.1 External evidence

Evidence-based methods compare outputs with supplied or retrieved knowledge. QAFactEval asks questions about summary content and checks answer recovery from the source; question generation and answerability affect its decisions (Fabbri et al., 2022). AlignScore learns alignment across heterogeneous tasks and aggregates source-chunk/output-sentence scores (Zha et al., 2023). Both support black-box generators but require trained auxiliary models. Question decomposition offers interpretable checks; learned alignment simplifies scoring. Neither eliminates verifier errors.

FactScore decomposes long answers into atomic facts and estimates their supported fraction relative to a designated source (Min et al., 2023). VeriScore instead restricts extraction to verifiable claims and uses surrounding context to resolve references (Song et al., 2024). Extraction therefore changes the score denominator, while missing retrieved support differs from contradiction. PrefixNLI specializes entailment to incomplete prefixes, enabling feedback during generation at the cost of verification calls (Harary et al., 2026). Detection becomes a decoding component, but remains conditional on verifier quality.

3.2 Sampling and consistency

SelfCheckGPT compares sentences with additional responses to the same prompt (Manakul et al., 2023). It avoids internal-access requirements but pays for repeated generation and consistency estimation. Agreement is only a proxy: samples can repeat a misconception, and comparison models can overlook disagreement.

Semantic uncertainty groups mutually entailing answers before estimating entropy, reducing variation caused by paraphrases (Kuhn et al., 2023; Farquhar et al., 2024). The weighted estimator uses probabilities; a discrete approximation uses cluster frequencies. These methods particularly target unstable confabulations, whereas stable false




Figure 2: The structure of this survey. Green leaves locate representative studies. Branches navigate detection signals, evaluation targets, and interventions, rather than define mutually exclusive classes. Tree design adapted from Liu et al. (2025); content is specific to this survey.

beliefs need not produce high entropy. Sampling temperature, sample count, and equivalence judgments affect the estimate. Consistency can triage responses, but cannot replace independent evidence when justification is required.

3.3 Internal signals

FOCUS combines token probabilities with attention-based uncertainty propagation from an accessible proxy model (Zhang et al., 2023b). The target generator can remain black-box; the proxy must expose both signals. Hidden-state classifiers learn truth labels from activations (Azaria and Mitchell, 2023). Lookback Lens uses context-versus-generation attention features with a supervised linear classifier (Chuang et al., 2024a). These approaches can avoid repeated responses but require the specified model signals and, for probes, representative labels.

A separable activation pattern need not represent truth itself. In controlled factual-recall experiments, Cheang et al. (2026) find that errors driven by learned associations overlap with correct outputs more than errors lacking those associations do. This qualifies claims that models internally recognize every falsehood. Evaluation must sepa-

rate knowledge familiarity, source adherence, and correctness, and test transfer to natural errors and different generators.

These families differ in what could falsify a decision. Alignment can expose a source contradiction but not establish source truth. Sampling tests claim reproducibility; probes test resemblance to labeled examples. Neither supplies independent factual evidence. Access requirements must distinguish the target generator, proxy, and verifier. Thus, “black-box” and “reference-free” are not interchangeable deployment labels; evidence, signals, supervision, and inference budgets need separate specification.

4 Benchmarks and Evaluation

Figure 3 separates answer quality from detector performance. Table 1 instantiates this distinction; its sizes count different objects, not relative difficulty. Reference frame, annotation unit, and sampling process precede metric selection.

4.1 Generation evaluation

TruthfulQA evaluates misconception-oriented answers for truthfulness and informativeness, with a separate multiple-choice setting (Lin et al., 2022).




Figure 3: Evaluation targets: answer quality versus error-label prediction. Shared corpora require distinct protocols; listed metrics are examples, not interchangeable scores. Tree layout follows Liu et al. (2024a); categorization is specific to text LLMs.

Table 1: Representative hallucination evaluation benchmarks. Dis: error-label discrimination; Gen: generated-content evaluation; MC: multiple choice. Sizes count different units. The three-rule design follows Liu et al. (2024a). Corpus and metric caveats appear in Appendix B.

Benchmark	Setting	Size	Focus	Metrics
TruthfulQA (Lin et al., 2022)	Gen/MC	817 questions	World factuality	Truthfulness; informativeness
HaluEval (Li et al., 2023)	Dis	35,000 samples	Mixed-task errors	Recognition accuracy
ANAH (Ji et al., 2024)	Dis	~12k sentences	QA support; contradiction	Type accuracy; F1
RAGTruth (Niu et al., 2024)	Dis	17,790 responses	RAG faithfulness	Response/span P, R, F1
ExpertQA (Malaviya et al., 2024)	Gen	2,177 questions	Correctness; attribution	Human factuality/support
TofuEval (Tang et al., 2024)	Dis/Gen	1,479 summaries*	Dialogue summaries	Balanced accuracy; error rate
HALoGEN (Ravichander et al., 2025)	Gen	10,923 prompts	Domain factuality	Atomic hallucination rate
FaithBench (Bao et al., 2025)	Dis/Gen	750 summaries	Summary faithfulness	Balanced accuracy; macro F1; conditional error rate
FactBench (Fatahi Bayat et al., 2025)	Gen	1,000 prompts	Verifiable claims	Hallucination score; precision

* Paper-reported retained count; the source arithmetic is inconsistent. HALoGEN’s count includes programming. FaithBench’s selected cases are not a prevalence sample. Dis/Gen indicates distinct possible uses, not a shared metric.

ExpertQA distinguishes correctness from attribution through expert questions and claim-level judgments (Malaviya et al., 2024). HALoGEN uses domain-specific verifiers; FactBench combines user-derived prompts with web verification (Ravichander et al., 2025; Fatahi Bayat et al., 2025). Their scores are not interchangeable because evidence and verifier scopes differ.

Source consistency likewise requires a utility measure. TofuEval studies topic-focused dialogue summaries, where factuality must coexist with relevance and completeness (Tang et al., 2024). Shorter answers may omit both errors and required information. Precision should accompany coverage and refusal behavior, with explicit empty/unverifiable-answer policies. Full benchmark sizes may also exceed this survey’s task scope: HALoGEN in-

cludes programming.

For claim-verification pipelines, comparisons should align extracted claims to annotations before scoring. Otherwise, differences between FActScore and VeriScore may reflect claim selection rather than verification quality (Min et al., 2023; Song et al., 2024). ExpertQA offers correctness and attribution judgments for this audit, whereas FactBench retains undecidable units (Malaviya et al., 2024; Fatahi Bayat et al., 2025).

4.2 Detector evaluation

HaluEval mixes generated task-specific examples with human-annotated general queries; recognition accuracy need not transfer to unselected errors (Li et al., 2023). ANAH annotates sentence-level types, evidence, and corrections; RAGTruth adds natu-

rally generated RAG responses with error spans (Ji et al., 2024; Niu et al., 2024). Response classification and span localization are different tasks: finding one error does not locate every unsupported claim.

Table 1 supplies labels, but not necessarily every detector’s inputs. RAGTruth supports response and span testing; attention-based detection additionally requires compatible generator traces, not merely stored source–response text (Niu et al., 2024). Regeneration to collect those signals changes the protocol and must be reported. Likewise, converting sentence scores into token labels adds a localization rule, not an original detector capability. TofuEval’s sentence and summary decisions require separate assessment (Tang et al., 2024). These distinctions connect method access to valid benchmark use.

FaithBench selects difficult cases through detector disagreement, making it a stress test rather than a prevalence estimate (Bao et al., 2025). TRIVIA+ examines long contexts and noisy supervision (Chen et al., 2026). Comparisons require matched splits, labels, generators, and evidence access, with thresholds fixed independently of test labels. Similar judge accuracy can conceal different error patterns and biased system-level estimates (Godbole and Jia, 2025). Mitigation must improve useful generation under an independently audited protocol.

5 Mitigation: Interventions and Assumptions

Figure 4 connects interventions to potential failure points; hybrid methods span stages. Faithfulness requires appropriate evidence use, whereas factuality also requires correct evidence. Evaluation must therefore respect the reference frame in Section 2.

5.1 Training and alignment

FaithDial supplies source-grounded dialogue revisions for generation and critic supervision (Dziri et al., 2022a). This requires curated examples. Self-Alignment instead derives preferences from a model’s self-evaluation, improves that evaluation through Self-Knowledge Tuning, and applies direct preference optimization (Zhang et al., 2024b). It reduces human labeling but depends on self-evaluation producing informative factuality preferences.

FactAlign uses sentence-level fine-grained Kahneman–Tversky optimization for credit assign-

ment within mixed-quality answers (Huang and Chen, 2024). UAlign incorporates confidence and semantic entropy into a reward model before proximal policy optimization (PPO) alignment (Xue et al., 2025). RHIO targets contextual faithfulness: masking retrieval heads produces negative examples, control tokens distinguish faithful from unfaithful continuations, and contrastive decoding subsequently separates them (Huang et al., 2025b). RHIO thus spans data construction, training, and inference. Labeling and optimization incur offline costs, while imperfect supervision can reinforce evaluator errors. Factuality rewards therefore require independent auditing.

5.2 Retrieval and prompting

FLARE retrieves actively when provisional continuations signal uncertainty, while context-faithful prompting changes instructions without retraining (Jiang et al., 2023; Zhou et al., 2023). Retrieval adds information but introduces selection errors and evidence conflicts. Self-RAG trains retrieval and reflection tokens that control when to retrieve and assess relevance, support, and utility; candidate continuations can then be ranked using these signals (Asai et al., 2024). This requires training and reflection supervision. Astute RAG instead uses a training-free, source-aware consolidation procedure: it elicits internal knowledge, compares it with retrieved passages, and iteratively resolves conflicts (Wang et al., 2025). This permits black-box use but adds calls and fallible reliability judgments. Comparison should separate retrieval quality, conflict resolution, and evidence use under matched corpora and budgets.

5.3 Decoding and internal intervention

Context-aware decoding (CAD) amplifies the difference between context-conditioned and context-free token distributions (Shi et al., 2024). For context c , query x , prefix $y_{<t}$, and contrast strength $\alpha \geq 0$,

$$p_{\text{CAD}}(y_t) \propto \frac{p_{\theta}(y_t \mid c, x, y_{<t})^{1+\alpha}}{p_{\theta}(y_t \mid x, y_{<t})^{\alpha}}. \quad (1)$$

CAD requires logits and additional decoding work. It strengthens context adherence without assessing context truth. DoLa contrasts mature-layer and dynamically selected earlier-layer predictions within the same model (Chuang et al., 2024b). DoLa needs layer outputs rather than an




Figure 4: Potential failure points and interventions, synthesized rather than causally validated. Hybrid methods span stages. Improvements require independent evaluation through Figure 3 and Table 1. Layout adapted from Liu et al. (2024a).

external corpus, exposing parametric knowledge where CAD favors contextual evidence.

TruthX learns separate semantic and truth-related latent spaces and edits internal representations along a learned direction (Zhang et al., 2024a). It requires learning auxiliary components. Lookback Lens trains a classifier over context-versus-generation attention ratios and selects candidate continuations with that classifier (Chuang et al., 2024a). Its mitigation closes the detector-generator loop, but also inherits detector errors. These mechanisms are unsuitable for text-only APIs without the required internal signals. Successful steering does not establish model-independent truth representations; transfer across architectures and natural errors needs testing.

5.4 Revision and abstention

Chain-of-Verification drafts an answer, generates verification questions, answers them independently, and revises the draft (Dhuliawala et al., 2024). Separating verification from the draft reduces answer contamination, but all stages can still share a parametric misconception. CRITIC adds tool feedback before revision, providing an external correction signal at the cost of tool calls and iterative generation (Gou et al., 2024). WikiChat also grounds dialogue through retrieval and verification, showing that these operations can be coordinated as a pipeline (Semnani et al., 2023). External feedback remains limited by search and tool reliability. Re-

vision must preserve supported information while removing errors.

Abstention can reduce unsupported assertions. SEAL learns a rejection token during training and uses its probability to regularize decoding (Huang et al., 2025a); it is more specific than appending a refusal instruction. Lower error can reflect fewer answers. Refusal and revision therefore need coverage and information retention alongside precision.

6 Analysis and Future Directions

Reference frames determine whether an intervention succeeds. A response can follow an incorrect document faithfully or contradict a document while being factually correct. Consequently, gains from CAD and attention-guided decoding answer a different question from gains on world-knowledge QA. We propose evaluating both source support and world correctness on the same examples, retaining an unknown label when evidence is insufficient. A controlled retrieval experiment should vary document correctness and order while keeping questions and the generator fixed. Astute RAG motivates the conflict axis; Stable-RAG specifically addresses sensitivity to retrieval permutations (Wang et al., 2025; Zhang et al., 2026). Measuring both axes would distinguish stronger evidence use from more reliable evidence selection.

Confidence and internal signals are useful proxies, not truth certificates. Repeated agreement and confident internal activation can arise from a

stable false association. [Cheang et al. \(2026\)](#) distinguish associated from unassociated hallucinations and find that the former resemble factual recall in the internal representations they study. This constrains broad interpretations of successful probing, uncertainty alignment, or representation editing: a detector may identify missing knowledge more readily than confidently recalled misinformation. Future evaluations should partition naturally generated errors by these mechanisms, hold out generator families, and test whether external verification improves detection specifically on confident falsehoods. Average accuracy conceals which failure process a detector recognizes.

Metric improvements must preserve useful information. Factual precision does not reward answering the difficult parts of a question. Controlled long-form experiments show that response length affects measured precision ([Zhao et al., 2025](#)), while ICAT explicitly measures the coverage of expected information aspects ([Samarinas et al., 2025](#)). Report unique supported claims, aspect coverage, and refusal alongside precision. Risk–coverage curves compare error rates at matched answered-question fractions. Revision additionally requires counts of corrected and newly introduced errors and retained supported claims; matched-length analysis tests whether gains mainly reflect less information.

Evaluation must remain independent of optimization. [Godbole and Jia \(2025\)](#) show that similar attribution-evaluator accuracy can conceal different false-positive/false-negative trade-offs and biased system-level error estimates. Optimizing and evaluating with the same verifier therefore risks rewarding its preferences. Independent auditing should report class-specific errors and test held-out domains and generators. Outdated reference facts also distort evaluation ([Jiang et al., 2026](#)), motivating versioned evidence and periodic label review. These requirements link Figure 4 back to Figure 3 and Table 1: improvements must survive independent checking under matched evidence, coverage, and budgets.

Resource trade-offs depend on the intervention stage. FactAlign pays for factuality feedback and parameter updates; Self-RAG combines offline supervision with online retrieval and candidate selection ([Huang and Chen, 2024](#); [Asai et al., 2024](#)). CoVe and CRITIC instead add verification and re-

vision calls during use ([Dhuliawala et al., 2024](#); [Gou et al., 2024](#)). Training-free operation can still be costly; training expenditure may be amortized across queries but repeated after domain changes. Report both stages, including tokens, requests, latency, and memory.

Detection and mitigation require joint validation. An intervention can improve generation while making remaining errors harder to detect. [Cheng et al. \(2025\)](#) find that reasoning prompts can reduce hallucinations while obscuring detection cues in their studied settings. This motivates testing detector–intervention combinations rather than assuming component performance transfers unchanged. Shared models or verifiers may also preserve a misconception across apparently confirmatory stages; this is a hypothesis to test, not a universal claim about composition.

A controlled study should compare a baseline, two individual interventions, and their combination under matched questions, evidence, and budgets. Retrieval changes require separate retrieval-quality reporting. Cross these systems with fixed detectors, recording residual generation errors, false acceptance, false alarms, and retained supported claims. Independent auditing and matched answered-question coverage would distinguish complementary correction from reduced answering or evaluator preference. This tests Figure 4’s interventions through Figure 3’s targets.

Limits of the synthesis. The evidence favors English text tasks and ACL-family venues; heterogeneous settings preclude controlled cross-paper rankings. Appendices A–C document reading depth and comparison boundaries, and Appendix D specifies an unperformed evaluation protocol.

7 Conclusion

RQ1 requires separating world factuality from source faithfulness and declaring the claim unit and unknown-label policy. For **RQ2**, external evidence enables support checks, whereas sampling and internal signals estimate reliability under distinct access and supervision constraints. For **RQ3**, useful mitigation requires independently verified gains in supported information at comparable coverage and cost. Future work should test detector–intervention combinations under evidence conflicts and distribution shifts, linking mechanism choice to residual errors.

References

- Vaibhav Adlakha, Parishad BehnamGhader, Xing Han Lu, Nicholas Meade, and Siva Reddy. 2024. [Evaluating correctness and faithfulness of instruction-following models for question answering](#). *Transactions of the Association for Computational Linguistics*, 12:681–699.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. 2024. [Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection](#). In *The Twelfth International Conference on Learning Representations*.
- Amos Azaria and Tom Mitchell. 2023. [The internal state of an LLM knows when it’s lying](#). In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 967–976, Singapore. Association for Computational Linguistics.
- Yejin Bang, Ziwei Ji, Alan Schelten, Anthony Hartshorn, Tara Fowler, Cheng Zhang, Nicola Cancedda, and Pascale Fung. 2025. [HalluLens: LLM hallucination benchmark](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24128–24156, Vienna, Austria. Association for Computational Linguistics.
- Forrest Sheng Bao, Miaoran Li, Renyi Qu, Ge Luo, Erana Wan, Yujia Tang, Weisi Fan, Manveer Singh Tamber, Suleman Kazi, Vivek Sourabh, Mike Qi, Ruixuan Tu, Chenyu Xu, Matthew Gonzales, Ofer Mendelevitch, and Amin Ahmad. 2025. [FaithBench: A diverse hallucination benchmark for summarization by Modern LLMs](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 448–461, Albuquerque, New Mexico. Association for Computational Linguistics.
- Baolong Bi, Shaohan Huang, Yiwei Wang, Tianchi Yang, Zihan Zhang, Haizhen Huang, Lingrui Mei, Junfeng Fang, Zehao Li, Furu Wei, Weiwei Deng, Feng Sun, Qi Zhang, and Shenghua Liu. 2025. [Context-DPO: Aligning language models for context-faithfulness](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 10280–10300, Vienna, Austria. Association for Computational Linguistics.
- Meng Cao, Yue Dong, and Jackie Cheung. 2022. [Hallucinated but factual! inspecting the factuality of hallucinations in abstractive summarization](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3340–3354, Dublin, Ireland. Association for Computational Linguistics.
- Arie Cattan, Paul Roit, Shiyue Zhang, David Wan, Roei Aharoni, Idan Szpektor, Mohit Bansal, and Ido Dagan. 2026. [Localizing factual inconsistencies in attributable text generation](#). *Transactions of the Association for Computational Linguistics*, 14:100–123.
- Chi Seng Cheang, Hou Pong Chan, Wenxuan Zhang, and Yang Deng. 2026. [Do LLMs really know what they don’t know? internal states mainly reflect knowledge recall rather than truthfulness](#). In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 713–730, San Diego, California, United States. Association for Computational Linguistics.
- Jiawei Chen, Hongyu Lin, Xianpei Han, and Le Sun. 2024. [Benchmarking Large Language Models in Retrieval-Augmented Generation](#). In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 17754–17762.
- Wenbo Chen, Veena Padmanabhan, Tootiya Giyahchi, Elaine Wong, and Leman Akoglu. 2026. [Rethinking evaluation for LLM hallucination detection: A desiderata, a new RAG-based benchmark, new insights](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14912–14931, San Diego, California, United States. Association for Computational Linguistics.
- Jiahao Cheng, Tiancheng Su, Jia Yuan, Guoxiu He, Jiawei Liu, Xinqi Tao, Jingwen Xie, and Huaxia Li. 2025. [Chain-of-thought prompting obscures hallucination cues in large language models: An empirical evaluation](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 1272–1305, Suzhou, China. Association for Computational Linguistics.
- Yung-Sung Chuang, Linlu Qiu, Cheng-Yu Hsieh, Ranjay Krishna, Yoon Kim, and James R. Glass. 2024a. [Lookback lens: Detecting and mitigating contextual hallucinations in large language models using only attention maps](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 1419–1436, Miami, Florida, USA. Association for Computational Linguistics.
- Yung-Sung Chuang, Yujia Xie, Hongyin Luo, Yoon Kim, James Glass, and Pengcheng He. 2024b. [DoLa: Decoding by Contrasting Layers Improves Factuality in Large Language Models](#). In *The Twelfth International Conference on Learning Representations*.
- Roi Cohen, May Hamri, Mor Geva, and Amir Globerson. 2023. [LM vs LM: Detecting factual errors via cross examination](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12621–12640, Singapore. Association for Computational Linguistics.
- Shehzaad Dhuliawala, Mojtaba Komeili, Jing Xu, Roberta Raileanu, Xian Li, Asli Celikyilmaz, and Jason Weston. 2024. [Chain-of-verification reduces hallucination in large language models](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 3563–3578, Bangkok, Thailand. Association for Computational Linguistics.
- Nouha Dziri, Ehsan Kamalloo, Sivan Milton, Omar Zaiane, Mo Yu, Edoardo M. Ponti, and Siva

- Reddy. 2022a. [FaithDial: A faithful benchmark for information-seeking dialogue](#). *Transactions of the Association for Computational Linguistics*, 10:1473–1490.
- Nouha Dziri, Hannah Rashkin, Tal Linzen, and David Reitter. 2022b. [Evaluating attribution in dialogue systems: The BEGIN benchmark](#). *Transactions of the Association for Computational Linguistics*, 10:1066–1083.
- Alexander Fabbri, Chien-Sheng Wu, Wenhao Liu, and Caiming Xiong. 2022. [QAFactEval: Improved QA-based factual consistency evaluation for summarization](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2587–2601, Seattle, United States. Association for Computational Linguistics.
- Cunhang Fan, Jun Zhang, Xue Zhang, Shuai Zhang, Zhao Lv, Jianhua Tao, and Zhengqi Wen. 2026. [ReFL: Reflective feedback learning for hallucination detection of large language models](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 19648–19665, San Diego, California, United States. Association for Computational Linguistics.
- Sebastian Farquhar, Jannik Kossen, Lorenz Kuhn, and Yarin Gal. 2024. [Detecting hallucinations in large language models using semantic entropy](#). *Nature*, 630(8017):625–630.
- Farima Fatahi Bayat, Lechen Zhang, Sheza Munir, and Lu Wang. 2025. [FactBench: A dynamic benchmark for in-the-wild language model factuality evaluation](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 33090–33110, Vienna, Austria. Association for Computational Linguistics.
- Shangbin Feng, Vidhisha Balachandran, Yuyang Bai, and Yulia Tsvetkov. 2023. [FactKB: Generalizable factuality evaluation using language models enhanced with factual knowledge](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 933–952, Singapore. Association for Computational Linguistics.
- Shangbin Feng, Weijia Shi, Yike Wang, Wenxuan Ding, Vidhisha Balachandran, and Yulia Tsvetkov. 2024. [Don’t hallucinate, abstain: Identifying LLM knowledge gaps via multi-LLM collaboration](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 14664–14690, Bangkok, Thailand. Association for Computational Linguistics.
- Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. 2023. [Enabling large language models to generate text with citations](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 6465–6488, Singapore. Association for Computational Linguistics.
- Zorik Gekhman, Jonathan Herzig, Roei Aharoni, Chen Elkind, and Idan Szpektor. 2023. [TrueTeacher: Learning factual consistency evaluation with large language models](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2053–2070, Singapore. Association for Computational Linguistics.
- Zorik Gekhman, Gal Yona, Roei Aharoni, Matan Eyal, Amir Feder, Roi Reichart, and Jonathan Herzig. 2024. [Does fine-tuning LLMs on new knowledge encourage hallucinations?](#) In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 7765–7784, Miami, Florida, USA. Association for Computational Linguistics.
- Aryo Pradipta Gema, Chen Jin, Ahmed Abdulaal, Tom Diethe, Philip Alexander Teare, Beatrice Alex, Pasquale Minervini, and Amrutha Saseendran. 2025. [DeCoRe: Decoding by contrasting retrieval heads to mitigate hallucinations](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 10003–10039, Suzhou, China. Association for Computational Linguistics.
- Ameya Godbole and Robin Jia. 2025. [Verify with caution: The pitfalls of relying on imperfect factuality metrics](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 22889–22912, Vienna, Austria. Association for Computational Linguistics.
- Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Nan Duan, and Weizhu Chen. 2024. [CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing](#). In *The Twelfth International Conference on Learning Representations*.
- Sapir Harary, Eran Hirsch, Aviv Slobodkin, David Wan, Mohit Bansal, and Ido Dagan. 2026. [PrefixNLI: Detecting factual inconsistencies as soon as they arise](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1414–1433, San Diego, California, United States. Association for Computational Linguistics.
- Or Honovich, Roei Aharoni, Jonathan Herzig, Hagai Taitelbaum, Doron Kukliansy, Vered Cohen, Thomas Scialom, Idan Szpektor, Avinatan Hassidim, and Yossi Matias. 2022. [TRUE: Re-evaluating factual consistency evaluation](#). In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 3905–3920, Seattle, United States. Association for Computational Linguistics.
- Chao-Wei Huang and Yun-Nung Chen. 2024. [FactAlign: Long-form factuality alignment of large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 16363–16375, Miami, Florida, USA. Association for Computational Linguistics.

- Lei Huang, Xiaocheng Feng, Weitao Ma, Yuchun Fan, Xiachong Feng, Yuxuan Gu, Yangfan Ye, Liang Zhao, Weihong Zhong, Baoxin Wang, Dayong Wu, Guoping Hu, Lingpeng Kong, Tong Xiao, Ting Liu, and Bing Qin. 2025a. [Alleviating hallucinations from knowledge misalignment in large language models via selective abstention learning](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24564–24579, Vienna, Austria. Association for Computational Linguistics.
- Lei Huang, Xiaocheng Feng, Weitao Ma, Yuchun Fan, Xiachong Feng, Yangfan Ye, Weihong Zhong, Yuxuan Gu, Baoxin Wang, Dayong Wu, Guoping Hu, and Bing Qin. 2025b. [Improving contextual faithfulness of large language models via retrieval heads-induced optimization](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16896–16913, Vienna, Austria. Association for Computational Linguistics.
- Denis Janiak, Jakub Binkowski, Albert Sawczyn, Bogdan Gabrys, Ravid Shwartz-Ziv, and Tomasz Jan Kajdanowicz. 2025. [The illusion of progress: Re-evaluating hallucination detection in LLMs](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 34728–34745, Suzhou, China. Association for Computational Linguistics.
- Ziwei Ji, Yuzhe Gu, Wenwei Zhang, Chengqi Lyu, Dahua Lin, and Kai Chen. 2024. [ANAH: Analytical annotation of hallucinations in large language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8135–8158, Bangkok, Thailand. Association for Computational Linguistics.
- Che Jiang, Biqing Qi, Xiangyu Hong, Dayuan Fu, Yang Cheng, Fandong Meng, Mo Yu, Bowen Zhou, and Jie Zhou. 2024. [On large language models’ hallucination with regard to known facts](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 1041–1053, Mexico City, Mexico. Association for Computational Linguistics.
- Xunyi Jiang, Dingyi Chang, Julian McAuley, and Xin Xu. 2026. [When benchmarks age: Temporal misalignment through large language model factuality evaluation](#). In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 500–512, Rabat, Morocco. Association for Computational Linguistics.
- Zhengbao Jiang, Frank Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu, Yiming Yang, Jamie Callan, and Graham Neubig. 2023. [Active retrieval augmented generation](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7969–7992, Singapore. Association for Computational Linguistics.
- Zhengping Jiang, Jingyu Zhang, Nathaniel Weir, Seth Ebner, Miriam Wanner, Kate Sanders, Daniel Khashabi, Anqi Liu, and Benjamin Van Durme. 2025. [Core: Robust factual precision with informative subclaim identification](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 19833–19856, Vienna, Austria. Association for Computational Linguistics.
- Katie Kang, Eric Wallace, Claire Tomlin, Aviral Kumar, and Sergey Levine. 2025. [Unfamiliar finetuning examples control how language models hallucinate](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3600–3612, Albuquerque, New Mexico. Association for Computational Linguistics.
- Anant Khandelwal, Manish Gupta, and Puneet Agrawal. 2025. [CoCoA: Confidence- and context-aware adaptive decoding for resolving knowledge conflicts in large language models](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 6835–6855, Suzhou, China. Association for Computational Linguistics.
- Kalpesh Krishna, Erin Bransom, Bailey Kuehl, Mohit Iyyer, Pradeep Dasigi, Arman Cohan, and Kyle Lo. 2023. [LongEval: Guidelines for human evaluation of faithfulness in long-form summarization](#). In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 1650–1669, Dubrovnik, Croatia. Association for Computational Linguistics.
- Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. 2023. [Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation](#). In *The Eleventh International Conference on Learning Representations*.
- Joosung Lee, Cheonbok Park, Hwiyeol Jo, Jeonghoon Kim, Joonsuk Park, and Kang Min Yoo. 2026. [Enhancing hallucination detection via future context](#). In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 731–749, San Diego, California, United States. Association for Computational Linguistics.
- Junyi Li, Jie Chen, Ruiyang Ren, Xiaoxue Cheng, Xin Zhao, Jian-Yun Nie, and Ji-Rong Wen. 2024. [The dawn after the dark: An empirical study on factuality hallucination in large language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10879–10899, Bangkok, Thailand. Association for Computational Linguistics.
- Junyi Li, Xiaoxue Cheng, Xin Zhao, Jian-Yun Nie, and Ji-Rong Wen. 2023. [HaluEval: A large-scale hallucination evaluation benchmark for large language models](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 6449–6464, Singapore. Association for Computational Linguistics.

- Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. [TruthfulQA: Measuring how models mimic human falsehoods](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252, Dublin, Ireland. Association for Computational Linguistics.
- Hanchao Liu, Wenyuan Xue, Yifei Chen, Dapeng Chen, Xiutian Zhao, Ke Wang, Liping Hou, Rongjun Li, and Wei Peng. 2024a. [A survey on hallucination in large vision-language models](#). *arXiv preprint arXiv:2402.00253*. Design reference; supplied version v2.
- Hanmeng Liu, Zhizhang Fu, Mengru Ding, Ruoxi Ning, Chaoli Zhang, Xiaozhang Liu, and Yue Zhang. 2025. [Logical reasoning in large language models: A survey](#). *arXiv preprint arXiv:2502.09100*. Design reference; supplied version v1.
- Junteng Liu, Shiqi Chen, Yu Cheng, and Junxian He. 2024b. [On the universal truthfulness hyperplane inside LLMs](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 18199–18224, Miami, Florida, USA. Association for Computational Linguistics.
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024c. [Lost in the middle: How language models use long contexts](#). *Transactions of the Association for Computational Linguistics*, 12:157–173.
- Liang Ma, Shuyang Cao, Robert L Logan IV, Di Lu, Shihao Ran, Ke Zhang, Joel Tetreault, and Alejandro Jaimes. 2023. [BUMP: A benchmark of unfaithful minimal pairs for meta-evaluation of faithfulness metrics](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12788–12812, Toronto, Canada. Association for Computational Linguistics.
- Matéo Mahaut, Laura Aina, Paula Czarnowska, Momchil Hardalov, Thomas Müller, and Lluís Marquez. 2024. [Factual confidence of LLMs: on reliability and robustness of current estimators](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4554–4570, Bangkok, Thailand. Association for Computational Linguistics.
- Chaitanya Malaviya, Subin Lee, Sihao Chen, Elizabeth Sieber, Mark Yatskar, and Dan Roth. 2024. [ExpertQA: Expert-curated questions and attributed answers](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3025–3045, Mexico City, Mexico. Association for Computational Linguistics.
- Potsawee Manakul, Adian Liusie, and Mark Gales. 2023. [SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9004–9017, Singapore. Association for Computational Linguistics.
- Radu Marinescu, Debarun Bhattacharjya, Junkyu Lee, Tigran T. Tchirakian, Javier Carnerero-Cano, Yufang Hou, Elizabeth M. Daly, and Alessandra Pascale. 2025. [FactReasoner: A probabilistic approach to long-form factuality assessment for large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 14547–14577, Suzhou, China. Association for Computational Linguistics.
- Sewon Min, Kalpesh Krishna, Xinxin Lyu, Mike Lewis, Wen-tau Yih, Pang Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023. [FactScore: Fine-grained atomic evaluation of factual precision in long form text generation](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12076–12100, Singapore. Association for Computational Linguistics.
- Dor Muhlgay, Ori Ram, Inbal Magar, Yoav Levine, Nir Ratner, Yonatan Belinkov, Omri Abend, Kevin Leyton-Brown, Amnon Shashua, and Yoav Shoham. 2024. [Generating benchmarks for factuality evaluation of language models](#). In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 49–66, St. Julian’s, Malta. Association for Computational Linguistics.
- Cheng Niu, Yuanhao Wu, Juno Zhu, Siliang Xu, KaShun Shum, Randy Zhong, Juntong Song, and Tong Zhang. 2024. [RAGTruth: A hallucination corpus for developing trustworthy retrieval-augmented language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10862–10878, Bangkok, Thailand. Association for Computational Linguistics.
- Abhilasha Ravichander, Shruti Ghela, David Wadden, and Yejin Choi. 2025. [HALoGEN: Fantastic LLM hallucinations and where to find them](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1402–1425, Vienna, Austria. Association for Computational Linguistics.
- Chris Samarin, Alexander Krubner, Alireza Salemi, Youngwoo Kim, and Hamed Zamani. 2025. [Beyond factual accuracy: Evaluating coverage of diverse factual information in long-form text generation](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 13468–13482, Vienna, Austria. Association for Computational Linguistics.
- Jakob Schuster, Vagrant Gautam, and Katja Markert. 2026. [Whose facts win? LLM source preferences under knowledge conflicts](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 29430–29459, San Diego, California, United States. Association for Computational Linguistics.

- Alessandro Scirè, Karim Ghonim, and Roberto Navigli. 2024. [FENICE: Factuality evaluation of summarization based on natural language inference and claim extraction](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 14148–14161, Bangkok, Thailand. Association for Computational Linguistics.
- Sina Semnani, Violet Yao, Heidi Zhang, and Monica Lam. 2023. [WikiChat: Stopping the hallucination of large language model chatbots by few-shot grounding on Wikipedia](#). In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 2387–2413, Singapore. Association for Computational Linguistics.
- Giovanni Servadio, Alessandro De Bellis, Dario Di Palma, Vito Walter Anelli, and Tommaso Di Noia. 2025. [Are the hidden states hiding something? testing the limits of factuality-encoding capabilities in LLMs](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6089–6104, Vienna, Austria. Association for Computational Linguistics.
- Weijia Shi, Xiaochuang Han, Mike Lewis, Yulia Tsvetkov, Luke Zettlemoyer, and Wen-tau Yih. 2024. [Trusting your evidence: Hallucinate less with context-aware decoding](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 783–791, Mexico City, Mexico. Association for Computational Linguistics.
- Shuzheng Si, Qingyi Wang, Haozhe Zhao, Yuzhuo Bai, Guanqiao Chen, Kangyang Luo, Gang Chen, Fan-chao Qi, Minjia Zhang, Baobao Chang, and Maosong Sun. 2026. [FaithLens: Detecting and explaining faithfulness hallucination](#). In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 14068–14099, San Diego, California, United States. Association for Computational Linguistics.
- Linxin Song, Taiwei Shi, and Jieyu Zhao. 2025. [The hallucination tax of reinforcement finetuning](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 2105–2120, Suzhou, China. Association for Computational Linguistics.
- Yixiao Song, Yekyung Kim, and Mohit Iyyer. 2024. [VeriScore: Evaluating the factuality of verifiable claims in long-form text generation](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 9447–9474, Miami, Florida, USA. Association for Computational Linguistics.
- Liyan Tang, Tanya Goyal, Alex Fabbri, Philippe Laban, Jiacheng Xu, Semih Yavuz, Wojciech Kryscinski, Justin Rousseau, and Greg Durrett. 2023. [Understanding factual errors in summarization: Errors, summarizers, datasets, error detectors](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11626–11644, Toronto, Canada. Association for Computational Linguistics.
- Liyan Tang, Igor Shalymov, Amy Wong, Jon Burnsky, Jake Vincent, Yu’an Yang, Siffi Singh, Song Feng, Hwanjun Song, Hang Su, Lijia Sun, Yi Zhang, Saab Mansour, and Kathleen McKeown. 2024. [TofuEval: Evaluating hallucinations of LLMs on topic-focused dialogue summarization](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 4455–4480, Mexico City, Mexico. Association for Computational Linguistics.
- Fei Wang, Xingchen Wan, Ruoxi Sun, Jiefeng Chen, and Sercan O Arik. 2025. [Astute RAG: Overcoming imperfect retrieval augmentation and knowledge conflicts for large language models](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 30553–30571, Vienna, Austria. Association for Computational Linguistics.
- Yuxia Wang, Minghan Wang, Muhammad Arslan Manzoor, Fei Liu, Georgi Nenkov Georgiev, Rocktim Jyoti Das, and Preslav Nakov. 2024. [Factuality of large language models: A survey](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 19519–19529, Miami, Florida, USA. Association for Computational Linguistics.
- Di Wu, Jia-Chen Gu, Fan Yin, Nanyun Peng, and Kai-Wei Chang. 2024. [Synchronous faithfulness monitoring for trustworthy retrieval-augmented generation](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 9390–9406, Miami, Florida, USA. Association for Computational Linguistics.
- Wenhao Wu, Wei Li, Xinyan Xiao, Jiachen Liu, Sujian Li, and Yajuan Lyu. 2023. [WeCheck: Strong factual consistency checker via weakly supervised learning](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 307–321, Toronto, Canada. Association for Computational Linguistics.
- Boyang Xue, Fei Mi, Qi Zhu, Hongru Wang, Rui Wang, Sheng Wang, Erxin Yu, Xuming Hu, and Kam-Fai Wong. 2025. [UAlign: Leveraging uncertainty estimations for factuality alignment on large language models](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6002–6024, Vienna, Austria. Association for Computational Linguistics.
- Xiaodong Yu, Hao Cheng, Xiaodong Liu, Dan Roth, and Jianfeng Gao. 2024. [ReEval: Automatic hallucination evaluation for retrieval-augmented large language models via transferable adversarial attacks](#). In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 1333–1351, Mexico City, Mexico. Association for Computational Linguistics.
- Yuheng Zha, Yichi Yang, Ruichen Li, and Zhiting Hu. 2023. [AlignScore: Evaluating factual consistency](#)

- with a unified alignment function. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11328–11348, Toronto, Canada. Association for Computational Linguistics.
- Jiaxin Zhang, Zhuohang Li, Kamalika Das, Bradley Malin, and Sricharan Kumar. 2023a. [SAC³: Reliable hallucination detection in black-box language models via semantic-aware cross-check consistency](#). In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 15445–15458, Singapore. Association for Computational Linguistics.
- Qianchi Zhang, Hainan Zhang, Liang Pang, Hong-Wei Zheng, and Zhiming Zheng. 2026. [Stable-RAG: Mitigating retrieval-permutation-induced hallucinations in retrieval-augmented generation](#). In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 25907–25926, San Diego, California, United States. Association for Computational Linguistics.
- Qinggang Zhang, Zhishang Xiang, Yilin Xiao, Le Wang, Junhui Li, Xinrun Wang, and Jinsong Su. 2025. [FaithfulRAG: Fact-level conflict modeling for context-faithful retrieval-augmented generation](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 21863–21882, Vienna, Austria. Association for Computational Linguistics.
- Shaolei Zhang, Tian Yu, and Yang Feng. 2024a. [TruthX: Alleviating hallucinations by editing large language models in truthful space](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8908–8949, Bangkok, Thailand. Association for Computational Linguistics.
- Tianhang Zhang, Lin Qiu, Qipeng Guo, Cheng Deng, Yue Zhang, Zheng Zhang, Chenghu Zhou, Xinbing Wang, and Luoyi Fu. 2023b. [Enhancing uncertainty-based hallucination detection with stronger focus](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 915–932, Singapore. Association for Computational Linguistics.
- Xiaoying Zhang, Baolin Peng, Ye Tian, Jingyan Zhou, Lifeng Jin, Linfeng Song, Haitao Mi, and Helen Meng. 2024b. [Self-alignment for factuality: Mitigating hallucinations in LLMs via self-evaluation](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1946–1965, Bangkok, Thailand. Association for Computational Linguistics.
- James Xu Zhao, Jimmy Z.J. Liu, Bryan Hooi, and See-Kiong Ng. 2025. [How does response length affect long-form factuality](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 3102–3125, Vienna, Austria. Association for Computational Linguistics.
- Wenxuan Zhou, Sheng Zhang, Hoifung Poon, and Muhao Chen. 2023. [Context-faithful prompting for large language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 14544–14556, Singapore. Association for Computational Linguistics.

A Selection, Provenance, and Reading Depth

This is a targeted narrative survey with a metadata snapshot of 1 October 2026, not an exhaustive systematic review or a meta-analysis. Discovery used official ACL Anthology collections for ACL, EMNLP, NAACL, EACL, Findings, and TACL across 2022–2026. Twenty-five collections were available; five requested year/venue combinations returned unavailable responses and were logged as such. A broad title filter for hallucination, factuality, faithfulness, attribution, uncertainty, retrieval, and related terms returned 1,766 records. Three targeted Anthology additions and seven external mechanism-defining studies produced 1,776 candidate records.

Selection prioritized direct relevance to text factuality or source faithfulness, distinct mechanisms, benchmark diversity, and evidence that qualified common assumptions. The working library contains 87 studies: 6 from 2022, 20 from 2023, 29 from 2024, 22 from 2025, and 10 from 2026. Eighty have Anthology records; the seven external records were checked against primary publication metadata. Canonical published versions were preferred over duplicate preprints. DOI and normalized-title checks found no exact duplicates among selected records; this is not a guarantee about near-duplicates in the reserve set. Automatic title triage excluded 401 candidates, and 1,288 remain in reserve without individual abstract assessment. These counts describe the actual workflow, not full-text exclusion decisions.

All 87 selected titles and abstracts were screened. Initial preparation inspected selected passages of 15 downloaded representative papers; drafting added targeted primary-source checks for the detailed detection, evaluation, and mitigation claims. This does not mean every selected paper was read in full. The companion evidence ledgers record claim-specific sections, access assumptions, and interpretive limits. The broad supplementary catalogue below distinguishes mechanisms without reproducing performance rankings. No experiments, independent benchmark replications, or statistical pooling were performed.

The manuscript cites 86 research studies, including citations embedded in its figures, and two survey papers credited solely for visual design. One intervention study remains in the 87-paper discovery library but is not cited in the manuscript. Library inclusion and final citation are separate editorial decisions. Figures 1–4 use the author’s supplied revised PDFs; their hashes and provenance are preserved in the repository. The example interactions are synthetic illustrations. The other figures organize published work and should not be interpreted as newly measured results. Bibliographic labels are generated with the unmodified ACL BibTeX style and checked against the labels printed inside the figures.

B Benchmark Units and Protocol Caveats

Table 2 supplies the counting and interpretation details behind Table 1. A benchmark may support both a description of generated errors and evaluation of a detector, but those uses require different denominators and sampling assumptions. In particular, reporting error frequencies within deliberately selected difficult examples does not estimate a generator’s error prevalence on ordinary requests.

C Supplementary Comparative Evidence

Tables 3–6 extend the main synthesis with complementary studies. The final column states the inference boundary used in this survey; it is not an assertion that the original authors omitted every proposed test. These studies expand coverage of mechanisms and failure conditions rather than form a ranked list. Descriptive entries are grounded in the recorded primary metadata and selected source passages, with reading depth retained in the repository.

Table 2: Interpretation safeguards for the nine main-text benchmarks.

Benchmark	Unit, protocol, and limitation
TruthfulQA	The 817 questions target common misconceptions. Generative truthfulness/informativeness and multiple-choice answer selection are distinct settings. MC is not hallucination-detector classification.
HaluEval	The 35,000 examples combine task-specific generated examples with a human-annotated general-query portion. Artificial and naturally generated errors should be reported separately rather than treated as one deployment distribution.
ANAH	Approximately 12,000 annotated sentences provide error types, evidence, and corrections. Generative versus discriminative annotators describe detector implementations; their task remains prediction of annotation labels.
RAGTruth	The 17,790 responses come from 2,965 prompts and six generators, not 17,790 independent documents. Response classification and span localization have different matching rules and must be evaluated separately.
ExpertQA	The 2,177 expert-curated questions support factuality and attribution analysis. Correctness and cited-source support are separate judgments; exact evaluation subsets and answer-generation conditions matter.
TofuEval	Section 3.3 reports 1,500 initial summaries, removal of 23, and 1,479 retained summaries. Those values are arithmetically inconsistent. We retain the reported 1,479 rather than invent a corrected dataset size. Topic relevance and completeness are separate from factual consistency.
HALoGEN	The 10,923 prompts cover all nine domains, including programming. That is the complete-suite count, not a verified text-only subset size. Different domain verifiers do not constitute one universal human truth oracle.
FaithBench	The 750 summaries derive from 75 passages and ten models, selected using detector disagreement. This supports difficult-case detector analysis; error frequencies are conditional on selection, not population prevalence.
FactBench	The 1,000-prompt snapshot uses user-derived prompts and fine-grained verification. Supported, unsupported, and undecidable units require explicit aggregation rules. A dynamic benchmark needs a fixed version for reproducible comparisons.

Table 3: Complementary evaluation designs.

Study	Mechanism or evaluation design	Comparative implication and boundary
TRUE (Honovich et al., 2022)	Standardizes factual-consistency datasets across tasks and evaluates metrics at the example level.	Instance discrimination differs from system-level correlation. Combined existing tasks do not estimate real-world hallucination prevalence.
AggreFact analysis (Tang et al., 2023)	Stratifies summarization error annotations by generating model and error type; metric gains vary across strata.	Pooled scores can hide weaknesses on newer generators or rare error types.
BUMP (Ma et al., 2023)	Uses minimally differing human-written summary pairs with one introduced error.	A controlled error-sensitivity probe complements natural generation data; minimal pairs do not reproduce the full distribution of model errors.
ALCE (Gao et al., 2023)	Evaluates answer quality and citation quality in end-to-end retrieval and cited generation.	Citation presence, citation support, and factual correctness are distinct criteria. A plausible citation is not by itself evidence.
LongEval (Krishna et al., 2023)	Studies fine-grained human judgments and partial annotation for long summaries.	Fine-grained annotation requires agreement and coverage reporting; partial annotation does not exhaustively verify all claims.
FACTOR (Muhlgay et al., 2024)	Transforms factual corpora into true versus false continuation choices.	Controlled factual choice evaluates a different task from open-ended generation and error localization.
ReEval (Yu et al., 2024)	Perturbs evidence to create dynamic tests of evidence use under retrieval.	Static accuracy may reflect memorized knowledge rather than source adherence; counterfactual evidence requires an explicit faithfulness target.
Core (Jiang et al., 2025)	Filters redundant or uninformative subclaims before factual precision aggregation.	Trivial repetitions can inflate precision; informativeness and coverage are complementary criteria.
TRIVIA+ (Chen et al., 2026)	Adds long-context RAG detection data and controlled sample-dependent label noise.	Long-context and noisy-label tests complement evaluation on clean labels.
RGB (Chen et al., 2024)	Separates noise robustness, negative rejection, information integration, and counterfactual robustness in English/Chinese RAG.	Retrieval conditions define distinct RAG abilities, not one interchangeable hallucination score.

Table 4: Complementary detection and explanation methods.

Study	Mechanism or evaluation design	Comparative implication and boundary
WeCheck (Wu et al., 2023)	Aggregates weak labels over actual generated samples, then learns a noise-aware classifier.	Supervision quality and error distribution distinguish a learned consistency model from zero-shot entailment.
FactKB (Feng et al., 2023)	Uses knowledge-base factual pretraining with entity/relation objectives for factuality evaluation.	Structured supervision targets relation errors; broader transfer requires separate evidence.
TrueTeacher (Gekhman et al., 2023)	Uses LLM-labeled model-generated summaries to train a smaller consistency evaluator.	Distillation trades annotation cost for teacher-label dependence; generated errors differ from artificial perturbations.
FENICE (Scirè et al., 2024)	Extracts atomic claims and aligns them to a source using NLI.	Local alignment improves interpretability, while claim extraction and entailment both remain possible error sources.
FactReasoner (Marinescu et al., 2025)	Models evidence–claim logical relationships probabilistically after decomposition/retrieval.	Joint reasoning differs from independently verifying claims; its result still depends on evidence retrieval and relation estimates.
Future context (Lee et al., 2026)	Samples continuations and uses them as clues for black-box hallucination detection.	Generated continuations extend consistency checking without supplying independent external verification.
FaithLens (Si et al., 2026)	Trains detection and explanations using filtered synthetic data, SFT, and rule-based RL.	Explanation quality and decision accuracy require separate validation; fluency alone does not establish explanation faithfulness.
QASemConsistency (Cattan et al., 2026)	Decomposes claims into predicate–argument QA pairs for fine-grained support localization.	Semantic error localization supplies a finer target than response-level binary classification.

Table 5: Evidence on reliability and failure mechanisms.

Study	Mechanism or evaluation design	Comparative implication and boundary
Factual confidence (Mahaut et al., 2024)	Compares confidence estimators and checks stability under meaning-preserving input changes.	Confidence can be unstable across paraphrases; access and supervision requirements differ.
New-knowledge fine-tuning (Gekhman et al., 2024)	Controlled closed-book QA experiments vary the amount of unfamiliar factual information in SFT.	The findings motivate knowledge-aware supervision within the studied fine-tuning setup.
Known facts (Jiang et al., 2024)	Examines layer-wise prediction dynamics when questions about the same fact yield correct versus incorrect answers.	Errors can reflect knowledge access rather than absence; measured patterns do not prove universal causal mechanisms.
Lost in the Middle (Liu et al., 2024c)	Evaluates how position of relevant evidence affects long-context use.	Evidence being present is insufficient for reliable use; its tasks are not themselves a comprehensive hallucination taxonomy.
Hidden-state limits (Servidio et al., 2025)	Tests factuality encoding on more realistic, model-dependent generated data.	Accuracy on synthetic true/false facts does not establish probe generalization to target-model outputs.
Unfamiliar supervision (Kang et al., 2025)	Relates hallucination response patterns to supervision of unfamiliar fine-tuning examples, including reward-model errors.	Factuality alignment depends on how unknowns and reward errors are labeled.
Illusion of Progress (Janiak et al., 2025)	Examines lexical-overlap labels and response-length heuristics in detector evaluation.	Semantic validation and simple baselines remain necessary; an LLM judge is not automatically a gold standard.
CoT obscures cues (Cheng et al., 2025)	Measures how reasoning prompts shift hidden/probability signals and detector effectiveness.	A generation intervention shifts the detector’s input distribution, making calibration transfer uncertain.
Hallucination tax (Song et al., 2025)	Studies RFT refusal degradation primarily on unanswerable mathematics, with reported transfer to factual QA.	Evidence is not a universal result for every factual-generation task.
Source preferences (Schuster et al., 2026)	Controlled synthetic-source study of source preferences and repetition in conflicting English evidence.	Frequency and source cues influence conflict resolution; synthetic-source preference is not proof of real-world source credibility.

Table 6: Complementary mitigation mechanisms.

Study	Mechanism or evaluation design	Comparative implication and boundary
FaithfulRAG (Zhang et al., 2025)	Models fact-level conflicts between retrieved and parametric knowledge before response generation.	Explicit reconciliation complements source preference; source adherence and world correctness remain distinct.
CoCoA (Khandelwal et al., 2025)	Adapts token-level decoding using confidence/context distribution differences.	Adaptive and fixed contrast differ in conflict sensitivity; low-conflict behavior also matters.
Context-DPO (Bi et al., 2025)	Constructs preferences from faithful versus stubborn responses under knowledge conflict.	A context-faithfulness objective differs from truthfulness with respect to external facts.

D A Common Reporting Protocol

The following protocol is a synthesis of the comparisons in Sections 3–6, not a benchmark run. It can be used to evaluate whether an intervention in Figure 4 solves the kind of error illustrated in Figure 1 under the target in Figure 3.

D.1 Fix the target before selecting a metric

Specify whether the target is world correctness, source support, or both; the answer unit; and the treatment of unknown or unverifiable claims. Record the generator/version, prompts, retrieval corpus/date, evidence truncation, decoding parameters, and judge/version. Use the same evidence and question distribution across systems where possible. If a method changes retrieval, separate retrieval quality from generation conditional on the retrieved passages.

For a nonempty extracted claim set C and fixed evidence E , a schematic support precision is

$$P_{\text{support}}(C, E) = \frac{\sum_{c \in C} \mathbf{1}[E \models c]}{|C|}. \quad (2)$$

This expression describes an aggregation principle, not the complete implementation of every factuality metric. Extraction determines the denominator; duplicate, trivial, subjective, and undecidable claims need explicit handling. An empty answer is not assigned perfect useful factuality by this definition. FActScore and VeriScore motivate claim-level analysis, while Core highlights the role of informative subclaims (Min et al., 2023; Song et al., 2024; Jiang et al., 2025).

D.2 Separate detection, localization, and generation quality

For binary detection, report class prevalence, confusion counts, precision/recall/F1, and balanced accuracy, $\frac{1}{2}(\text{TPR} + \text{TNR})$, where appropriate. Fix thresholds on independent validation data. For spans, specify tokenization and exact versus overlap matching. For continuous scores, report discrimination and calibration separately; neither establishes the truth of an individual answer. Inspect false-positive and false-negative examples across error types, evidence lengths, and generator families.

For generation, pair factuality with unique supported content and required-aspect coverage. A source-faithful answer may still omit the answer to the question. Human auditing should sample both high- and low-scoring outputs, use explicit instructions, and report disagreement or adjudication. LongEval provides relevant human-evaluation guidance (Krishna et al., 2023). Avoid pooling scalar scores whose claim extraction, evidence, or unknown-label policy differs.

D.3 Measure selection and revision effects

Let A_τ be the set of answers retained by a confidence threshold τ among N questions. For nonempty A_τ , define

$$\text{Coverage}(\tau) = \frac{|A_\tau|}{N}, \quad \text{Risk}(\tau) = \frac{\sum_{i \in A_\tau} \mathbf{1}[\text{answer } i \text{ is erroneous}]}{|A_\tau|}. \quad (3)$$

The error predicate must use the declared reference frame. Comparing risk at matched coverage prevents a system from appearing universally better merely by refusing more questions. This response-level selective-evaluation definition is distinct from SEAL’s token-level training objective. For revision, count corrected

errors, newly introduced errors, retained supported claims, and required information lost. Measure both before and after revision on the same initial drafts.

D.4 Audit independence, robustness, and cost

A verifier used for training or candidate selection should not be the sole final judge. Use independently collected labels or an independently audited evaluator. Test natural errors from held-out generators, difficult consistent falsehoods, evidence conflicts, context permutations, paraphrases, and updated facts. Disaggregate performance rather than interpreting an average as proof of transfer. Verify with Caution and temporal benchmark analysis motivate independent error inspection and timestamped evidence (Godbole and Jia, 2025; Jiang et al., 2026).

Report offline annotation and training costs separately from inference calls, sampled tokens, retrieval requests, tool invocations, latency, and memory. A black-box method can be expensive, while an internal method can be unavailable even if computationally economical. Do not infer a fixed latency multiplier from the number of model distributions or stages without measurement. Comparisons should either match budgets or present a quality–cost trade-off.

D.5 Method access and deployment trade-offs

Table 7: Mechanism-level deployment comparison. Costs are qualitative resources, not measured rankings.

Family / example	Required access or supervision	Principal extra resource	Important boundary
Source alignment / AlignScore	Source and output text; trained auxiliary model	Verifier inference over chunks	Auxiliary-model errors and evidence truncation
Claim verification / FActScore	Claim extraction, knowledge source, retrieval and verifier	Search and per-claim checks	Retrieval failure differs from disproof
Self-checking / Self-CheckGPT	Repeated outputs; no generator internals	Additional generation and comparisons	Repeated falsehood can appear consistent
Semantic entropy	Samples and semantic equivalence; likelihoods for weighted estimation	Sampling and clustering	Estimator access differs by variant
Probe / Lookback Lens	Attention maps and labeled training spans	Feature extraction and probe inference	Supervised contextual proxy, not proof of truth
Factuality alignment / FactAlign	Generator training and sentence-level factual feedback	Verification plus optimization	Reward errors can be reinforced
Retrieval control / Self-RAG	Training, retrieval, reflection-token supervision	Training and retrieved candidates	Reflection judgments need auditing
Contrastive decoding / CAD	Context-conditioned and context-free logits	Extra distribution at decoding	Greater source adherence can amplify false evidence
Layer contrast / DoLa	Intermediate layer predictions	Layer-output computation	Parametric knowledge is not external evidence
Representation editing / TruthX	Hidden states and learned auxiliary components	Auxiliary learning and intervention	Transfer beyond training conditions is uncertain
Verification / CoVe, CRITIC	Model calls; CRITIC also needs tools	Verification and revision rounds	Isolated prompts are not independent knowledge
Selective generation / SEAL	Rejection-token training and regularized decoding	Training; altered answering behavior	Lower risk must be interpreted at matched coverage